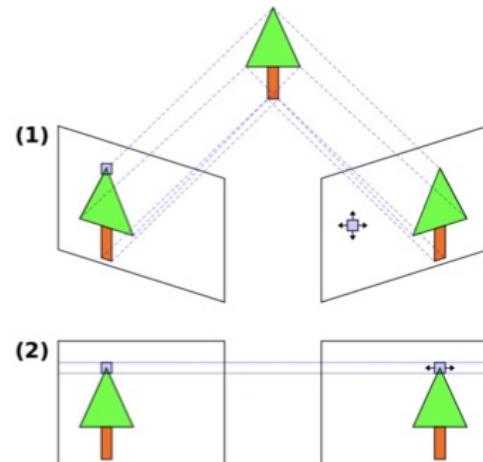
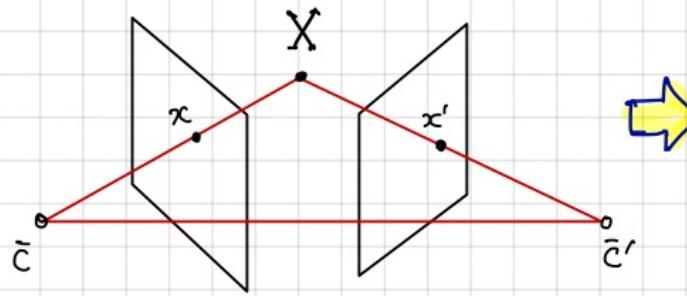


LECTURE 6 : STEREO + MRF

The goal of image rectification:

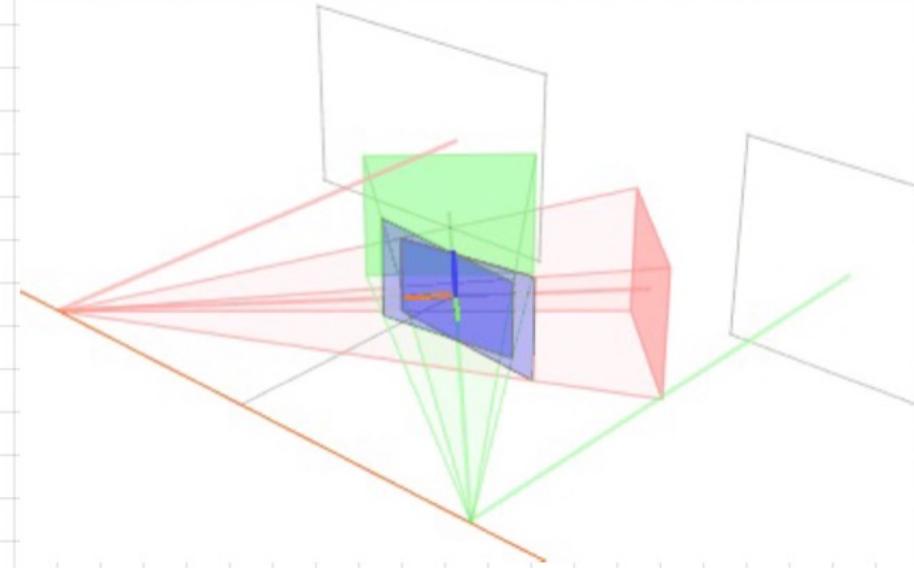


the search space before (1) and after (2) rectification



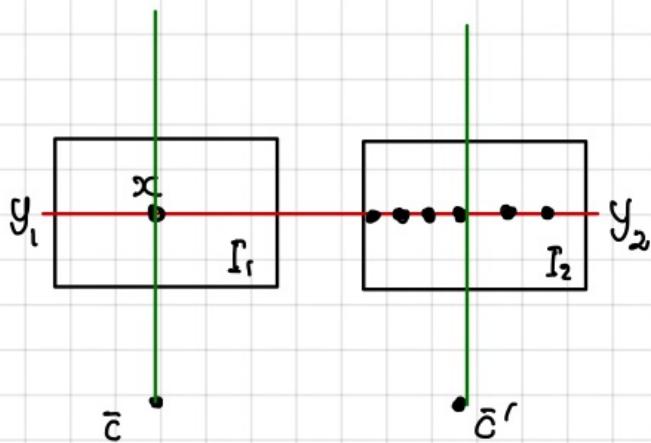
$$I_1(x) = I_2(x')$$

Image rectification process:



source: Wikipedia

[0] Disparity & Depth

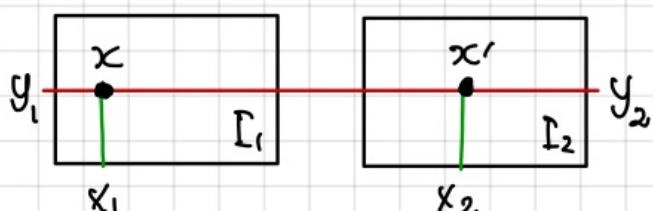


All the points on the red line on the right image are the candidates of a point, x' , that make $I_2(x') = I_1(x)$



where x & x' represents the same point in the world.

This red line is thus useful to find x' given x . Since, now we don't need to search x' in the entire image of I_2 .



$$I_1(x) = I_2(x')$$

$$I_1(x_1, y_1) = I_2(x_2, y_2)$$

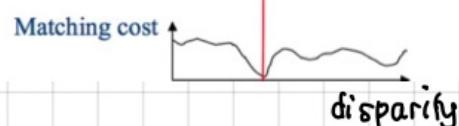
$$y_1 = y_2$$

$$\text{Disparity} \rightarrow d = |x_1 - x_2|$$

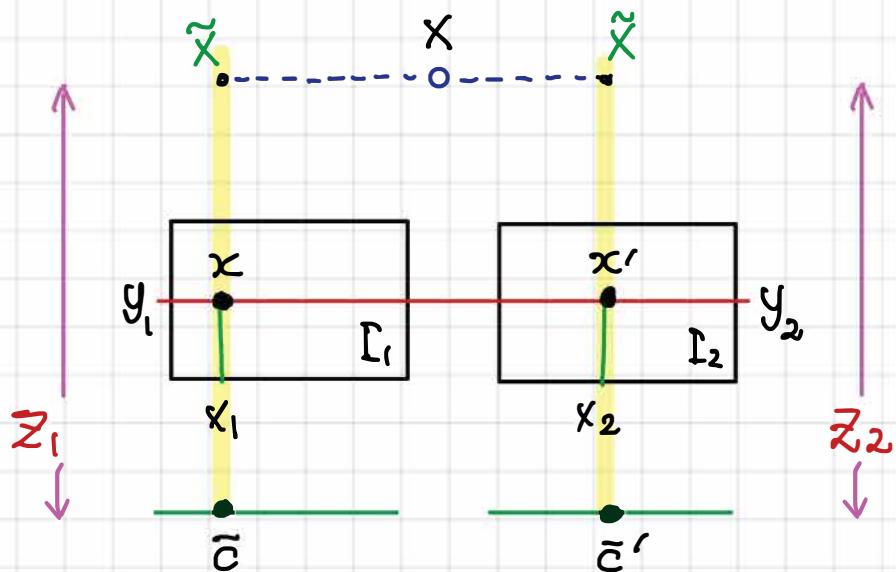
The larger the disparity, the smaller the depth: $d \propto 1/z$



left image

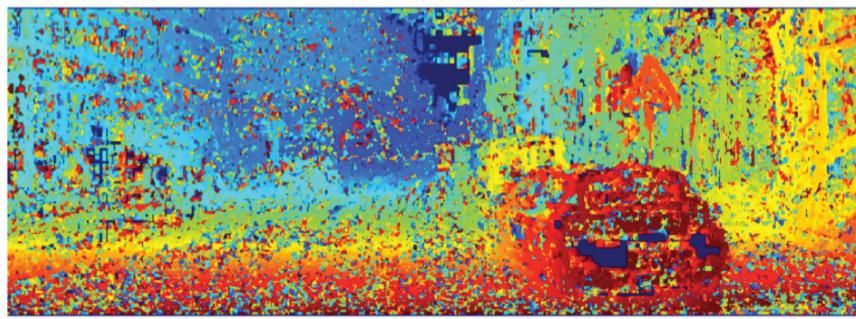


#6
For a pair of rectified images, there is only one disparity or depth map:

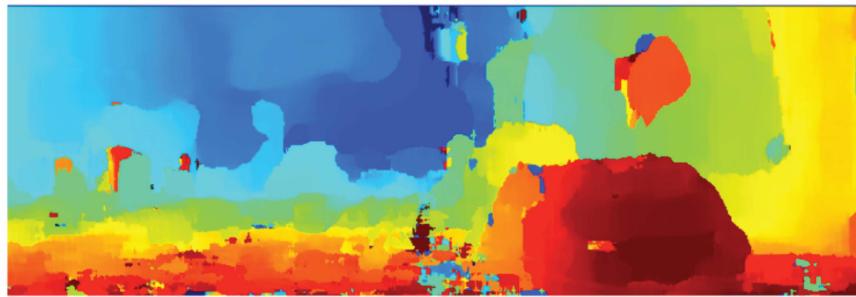


[3] Disparity Map from Rectified Images

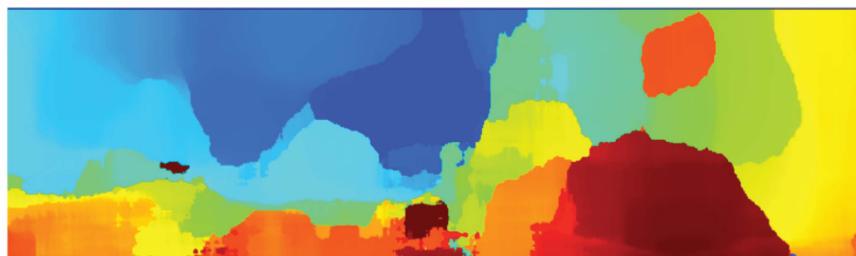
Pixel or patch based operations:



patch size = 5



patch size = 35



patch size = 85

Problem: small patches \rightarrow more details but more noise
large patch \rightarrow less noise but less details

Solution: MRF (Markov Random Fields)

[•] Markov Random Field (MRF)

Markov Random Field (MRF) is one type of graphical models



Motivation : Graphical Models

① An image is a set of numbers representing light intensity / brightness and arranged orderly in a 2D space.

② To make sense those numbers, we need to find the correlations among the numbers.

③ To find the correlations means to formulate the numbers spatially.

④ One of the methods to formulate is through the graphical models

A graphical model is a probabilistic model for which a graph expresses the conditional dependence structure between random variables.

This requires us to understand the basic principles of probabilistic inference.

Basic Probability :

Product rule

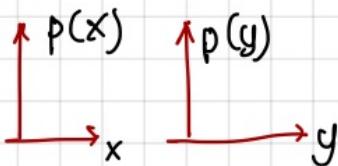
Independent variables

$$p(x|y) = p(x)$$

$$p(y|x) = p(y)$$

Thus:

$$p(x,y) = p(x)p(y)$$



Notes:

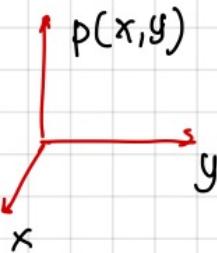
$$\textcircled{1} \quad 0 \leq p(x) \leq 1$$

$$\textcircled{2} \quad \int p(x) dx = 1$$

$$\sum_x p(x) = 1$$

Dependent variables

$$\begin{aligned} p(x,y) &= p(x|y)p(y) \\ &= p(y|x)p(x) \end{aligned}$$



Sum rule (Marginalization)

Two dependent variables x and y:

$$p(x) = \sum_y p(x,y)$$

Example:

x & y are random variables representing the possible faces of a coin

$$x, y \in \{\text{Head, Tail}\}$$

Then:

$$p(x=\text{H}) = \sum_{y \in \{\text{H, T}\}} p(x=\text{H}, y)$$

$$= \underbrace{p(x=\text{H}, y=\text{H})}_{1/4} + \underbrace{p(x=\text{H}, y=\text{T})}_{1/4}$$

↓

Three dependent variables, x, y, z:

$$p(x) = \sum_y \sum_z p(x,y,z)$$

$$p(x,y) = \sum_z p(x,y,z)$$

Bayes' Theorem:

$$p(x,y) = p(x|y) p(y) \quad (\text{by product rule})$$

$$p(x|y) = \frac{p(x,y)}{p(y)}$$

Thus :

$$p(x|y) = \frac{p(y|x) p(x)}{p(y)}$$

↑
Posterior = $\frac{\text{likelihood} \times \text{prior}}{\text{evidence}}$



$$p(x|d) = \frac{p(d|x) p(x)}{p(d)}$$

where :

d = the observed data (the value known to us)

x = the hidden / random variable (the value to estimate)

$p(x|d)$ = the probability of x having a certain value given the data, d .

$p(d|x)$ = the likelihood probability of the data assuming x having a certain value.

$p(x)$ = the prior probability of x

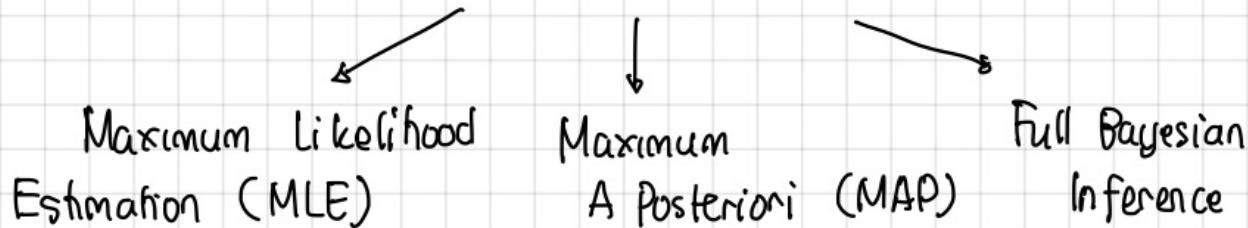
$p(d)$ = the evidence , where :

$$p(d) = \sum_x p(d,x) \quad (\text{by sum rule})$$

$$= \sum_x p(d|x) p(x) \rightarrow \text{expensive to compute if the dimension of } x \text{ is large.}$$

$$p(d) = \int p(d|x) p(x) dx \rightarrow \text{in most cases it's intractable to integrate.}$$

[•] Probabilistic Inference Methods



$$x^* = \underset{\{x\}}{\operatorname{argmax}} \underset{\text{likelihood}}{p(d|x)}$$

$$\begin{aligned} x^* &= \underset{\{x\}}{\operatorname{argmax}} p(x|d) \\ &= \underset{\text{likelihood}}{\operatorname{argmax}} p(d|x) \underset{\text{prior}}{p(x)} \end{aligned}$$

x^* = the most probable value of x .

$\{x\}$ = a set of possible values of x
(the candidates of x)

argmax = to return the value of x
that maximizes $p(d|x)$

$$p(x|d) = \frac{p(d|x) p(x)}{p(d)}$$

it's intractable to solve.

Example: According to our observation, a patient has a set of symptoms (= coughing, headache, fever), d .

We also have a set of possible illness: flu, infection and cancer.

The intuitive meaning of MLE:

1. Take each possible illness (e.g. $x = \text{flu}$)
2. Verify if it fits to the symptoms \rightarrow probability of x .
3. Decide the illness based on the highest probability

[•] Maximum A Posteriori (MAP)

$$x^* = \underset{\{x\}}{\operatorname{argmax}} p(x|d) = \underset{\{x\}}{\operatorname{argmax}} \frac{p(d|x) p(x)}{p(d)}$$

$$x^* = \underset{\{x\}}{\operatorname{argmax}} p(d|x) p(x) \quad \leftarrow \text{Because } p(d) \text{ is independent from } x.$$

Notes:

1. The weakness of MLE is that it allows impossible candidates, as it treats candidates uniformly.
2. MAP treats the candidates differently : $p(x=\text{flu}) > p(x=\text{cancer})$
3. Main criticism of MAP is how to determine $p(x)$.

[6] Full Bayesian Inference

$$p(x|d) = \frac{p(d|x) p(x)}{p(d)}$$

1. This is the ideal and most robust inference that besides providing a prediction on the value of x , it also provides the uncertainty, or how high/ low the probability of the prediction.
2. It's possible because of $p(d)$, the evidence. This will normalize the product of $p(d|x) p(x)$.
3. However to obtain $p(d)$ is analytically intractable, and in discrete, it can be prohibitively expensive.

[•] Graphical Models: Factorization

Graphical model: A probabilistic model where a graph expresses the conditional dependence structure between random variables.



Conditional Independence:

$$p(x_1 | x_2, x_3, x_4) = p(x_1 | x_2)$$



If implies that x_1 is independent of x_3 and x_4 given x_2 .



Factorization of joint probability

$$p(x_1, x_2, x_3, x_4) =$$

$$p(x_3 | x_2, x_1) \\ p(x_2 | x_1) p(x_1)$$

If x_3 is conditional independent of x_1 , then:

$$p(x_1, x_2, x_3) = p(x_3 | x_2) p(x_2 | x_1) \\ p(x_1)$$



Factorization reduces the complexity of computing a joint probability



Computing $p(x_1, x_2, x_3)$

means we calculate all possible values of

x_1, x_2, x_3 . E.g. $x_i \in \{0, 1\}$

then: $x_1 = 0, x_2 = 0, x_3 = 0$

$x_1 = 0, x_2 = 0, x_3 = 1$

$x_1 = 0, x_2 = 1, x_3 = 0$

\vdots

Graphical model is useful due to factorization.

If we use factorization:

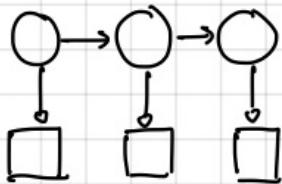
$$p(x_1, x_2, x_3) = p(x_1) p(x_2) p(x_3)$$

the computation is only: 2^n

[•] Type of Graphical Models

Directed Graphs:

e.g.:



Formulation:

$$p(x_1, \dots, x_N) = \prod_n p(x_n | x_{\text{pa}(x_n)})$$

where:

$x_{\text{pa}}(x_n)$ = the parents of x_n

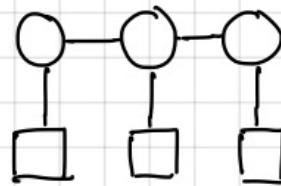
Example:



$$p(x_1, x_2, x_3) = p(x_3 | x_2) \\ p(x_2 | x_1) p(x_1)$$

Undirected Graphs

e.g.:



Formulation:

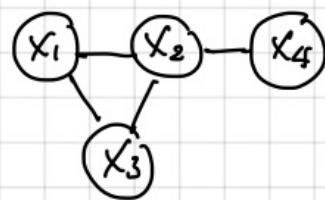
$$p(x_1, \dots, x_N) = \frac{1}{Z} \prod_c \phi_c(x_1, \dots, x_N)$$

where: C = a set of cliques.

Z = a normalization factor.

A clique is a set of mutually connected nodes.

e.g.:



Cliques: $\{\{x_1, x_2, x_3\}, \{x_2, x_4\}\}$



$$p(x_1, \dots, x_4) = \frac{\phi[x_1, x_2, x_3] \phi[x_2, x_4]}{\sum_{x_1} \sum_{x_2} \sum_{x_3} \sum_{x_4} \phi[x_1, x_2, x_3] \phi[x_2, x_4]}$$

$$Z = \sum_{x_1} \sum_{x_2} \sum_{x_3} \sum_{x_4} \phi[x_1, x_2, x_3] \phi[x_2, x_4]$$

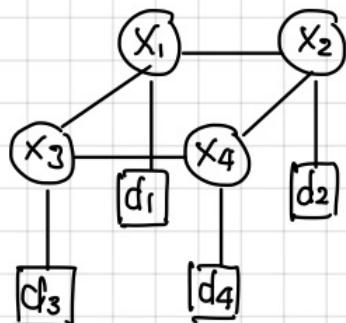
This Z is prohibitively expensive to compute.

Most undirected graphs utilize MAP to ignore Z .

[•] MRF : Formulation

A Markov random field (MRF) is a (loopy undirected graph).

e.g. :



Based on the cliques:

$$\begin{aligned}
 p(\{x\}, \{d\}) &= p(x_1 \dots x_4, d_1 \dots d_4) \\
 &= \frac{1}{Z} \phi[x_1, d_1] \phi[x_2, d_2] \phi[x_3, d_3] \phi[x_4, d_4] \\
 &\quad \phi[x_1, x_2] \phi[x_2, x_4] \phi[x_4, x_3] \phi[x_1, x_3] \\
 &= \frac{1}{Z} \prod_i^4 \phi[x_i, d_i] \prod_{j \in N_i} \phi[x_i, x_j]
 \end{aligned}$$

General MRF's formulation:

$$p(\{x\}, \{d\}) = \frac{1}{Z} \prod_i^N \phi[x_i, d_i] \prod_{j \in N_i} \phi[x_i, x_j]$$

Using MAP:

$$\{x\}^* = \underset{\{x\}}{\operatorname{argmax}} p(\{x\}, \{d\}) \propto \underset{\{x\}}{\operatorname{argmax}} p(\{x\} | \{d\})$$

$$= \underset{\{x\}}{\operatorname{argmax}} \frac{1}{Z} \prod_i \phi[x_i, d_i] \prod_{j \in N_i} \phi[x_i, x_j]$$

$$\propto \underset{\{x\}}{\operatorname{argmax}} \sum_i \left[\log \phi[x_i, d_i] + \sum_{j \in N_i} \log \phi[x_i, x_j] \right]$$

$$\{x\}^* = \underset{i}{\operatorname{argmin}} \sum_i f_d(x_i, d_i) + \sum_{j \in N_i} f_p(x_i, x_j)$$

[•] MRF : Data and Prior Terms

$$\{x\}^* = \operatorname{argmin}_i \sum_i \frac{f_d(x_i, d_i)}{\text{Data term}} + \sum_{j \in N_i} \frac{f_p(x_i, x_j)}{\text{Prior term}}$$

↙ ↘

Data term

Prior term:

The basic idea of determining the data term:

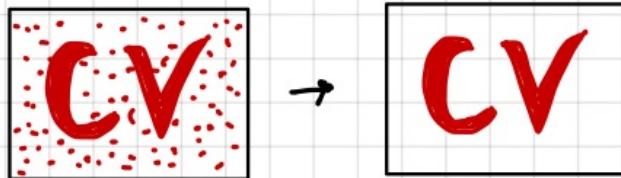
Due to argmin , $f_d(x_i, d_i)$ should produce a smaller value for a more probable x_i 's value w.r.t. d_i .

The basic idea of determining the prior term (smoothness prior):

$f_p(x_i, x_j)$ should generate a smaller value when x_i & x_j have the same value.

↓

Example: Denoising



1. $x_i \in \{0, 1\}$ or $x_i = \{B, F\}$ where $x_i \rightarrow \begin{cases} 0 \rightarrow \text{white}, \\ 1 \rightarrow \text{red} \end{cases}$
2. Prior term: $f(x_i, x_j) = |x_i - x_j|$: We encourage them to be the same.
3. Data term: The idea is that if d_i is white, x_i should be 0 or B. and if d_i is red, x_i should be 1 or F.

(1) Transform x_i 's value to d_i 's domain \rightarrow Because $x_i = \{0, 1\}$ and $d_i = \{0, \dots, 255\}$

$$T(x_i) = \begin{cases} \text{white} = (255, 255, 255) & \text{if } x_i = 0 \\ \text{red} = (255, 0, 0) & \text{if } x_i = 1 \end{cases}$$

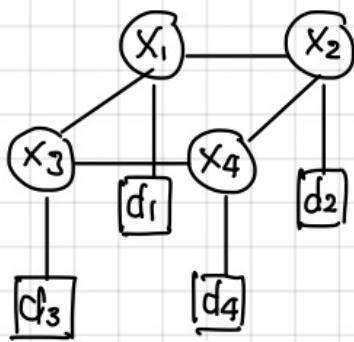
(2) Define a distance function between $T(x_i)$ and d_i .

$$\text{e.g.: } f_d(x_i, d_i) = f_d(T(x_i), d_i) = |T(x_i) - d_i|$$

Since $x_i = \{0, 1\}$, thus there are 2 possible values of f_d :

$$f_d(T(x_i=0), d_i) \text{ and } f_d(T(x_i=1), d_i).$$

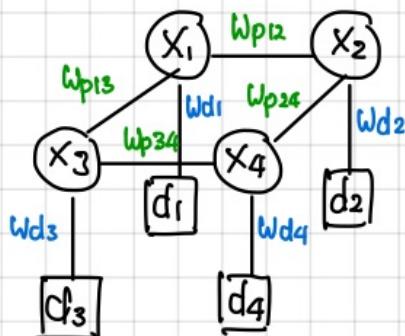
[•] MRF : Graph Optimization



$$\{x\}^* = \arg \min \sum_i f_d(x_i, d_i) + \sum_{i \in N_i} f_p(x_i, x_j)$$

↓

Having the values of f_d and f_p for every x_i means that we have a complete weighted graph.



Q : How from this weighted graph, can we estimate the optimum values of $\{x\}$?

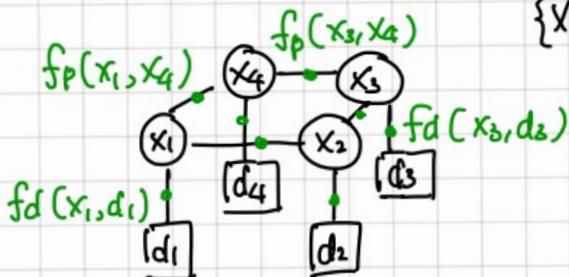
A : Graph optimization methods

Belief Propagation (BP) Graphcuts

[●] MRF : Optimization

[i] Motivation

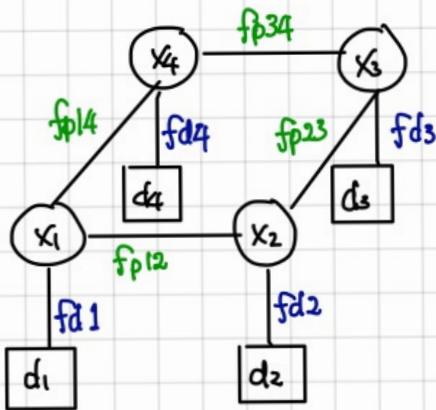
MRF :



$$\{x_1, \dots, x_4\}^* = \underset{\{x\}}{\operatorname{argmin}} \sum_i f_d(x_i, d_i) + \sum_{j \in N_i} f_p(x_i, x_j)$$

the value of each term indicates the values of the edges in the graph.

Having defined the data term (f_d) & the prior term (f_p) for every node (pixel), we can now have a complete graph :



Q : How can we optimize the graph so that we can predict the labels of x_1, x_2, x_3, x_4 ?

Q : What does it mean by 'optimizing' a graph ?

A : Example : $\textcircled{x}_1 \rightarrow \textcircled{x}_2$ & $x_i \in \{0, 1\}$

$$\{x_1, x_2\}^* = \underset{\{x_1, x_2\}}{\operatorname{argmin}} f_p(x_1, x_2)$$

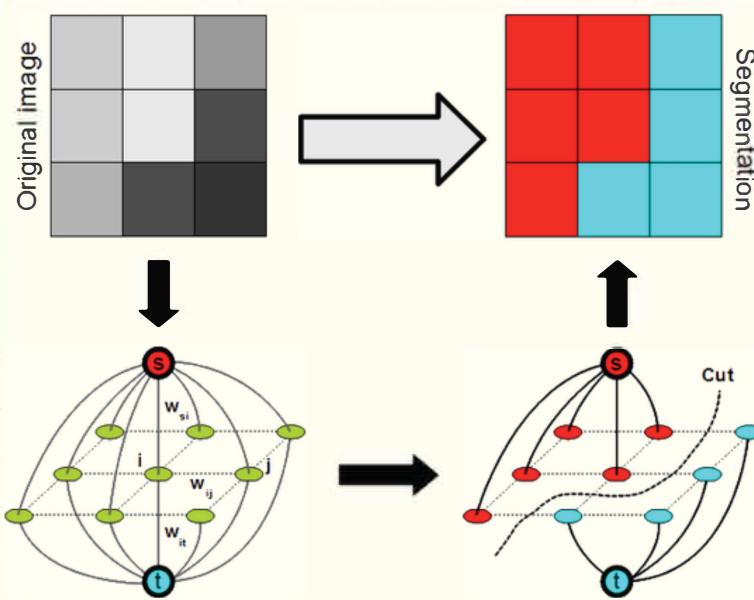
$$= \underset{\{x_1, x_2\}}{\operatorname{argmin}} \left[f_p(x_1=0, x_2=0), f_p(x_1=1, x_2=0), f_p(x_1=0, x_2=1), f_p(x_1=1, x_2=1) \right]$$

We find the values of x_1 & x_2 that minimize $f_p()$, or to choose one of four possible configurations.

[2] Graphcuts

There are two methods of graph optimization: Graphcuts & Belief Propagation.
Here, we focus on graphcuts, particularly on binary graphcuts.

Basic concept: Imagine we have a 3×3 image. Some of the pixels belong to the foreground & some belong to the background. However, due to noise/outlier and color variations, we are not sure & thus want to determine which pixels belongs to which label.



[3] Graphcuts : Data & Prior Terms

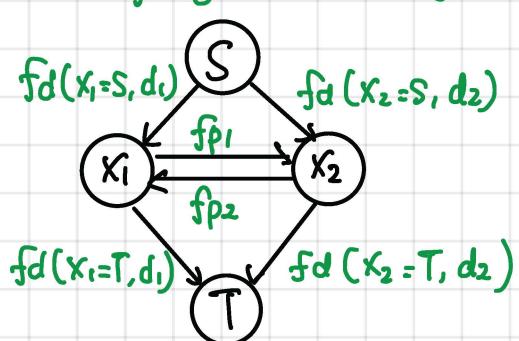
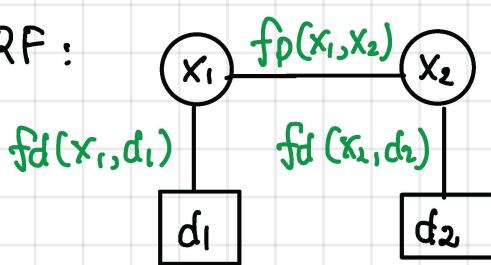
Q: How to assign the data & prior terms in a graph?

A: Suppose we have: x_1 & x_2 where $x_i = \{S, T\}$

and d_1, d_2

↑
foreground background

MRF:

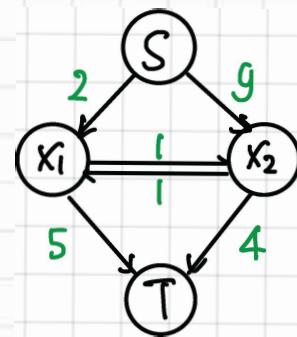
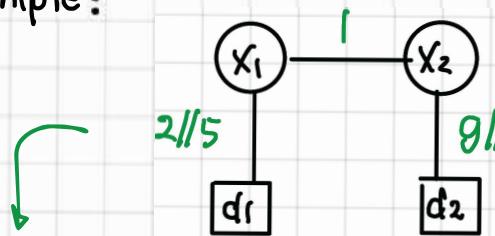


Assume:

$$f_p(x_1=T, x_2=S) = f_p(x_1=S, x_2=T) \rightarrow f_{p1} = f_{p2}$$

$$f_p(x_1=T, x_2=T) = f_p(x_1=S, x_2=S) = 0$$

Example:



$$fd(x_1 = S, d_1) = 2 \quad ; \quad fd(x_2 = S, d_2) = 9$$

$$fd(x_1 = T, d_1) = 5 \quad ; \quad fd(x_2 = T, d_2) = 4$$

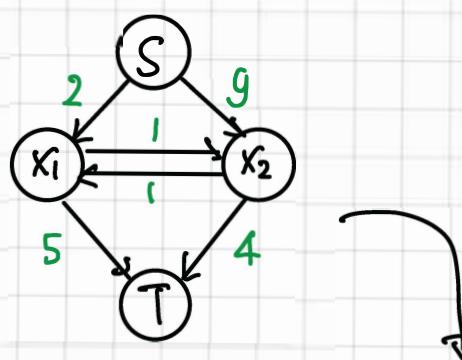
$$fp(x_1 = S, x_2 = T) = fp(x_1 = T, x_2 = S) = 1$$

$$fp(x_1 = S, x_2 = S) = fp(x_1 = T, x_2 = T) = 0$$

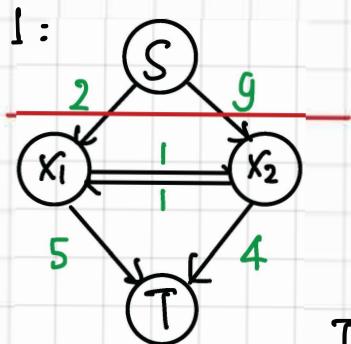
[4] Graphcuts: Max-Flow / Min-Cut

Q: How to optimize the graph, so that we know the labels of x_1 and x_2 ?

A:



cut 1:



$$\text{cut cost: } 9 + 2 = 11$$

↓

This cut implies:

This is for minimization. For maximization: $x_1 = T$
 $x_2 = T$

$$\left\{ \begin{array}{l} x_1 = S \\ x_2 = S \end{array} \right.$$

$$\text{cut cost: } 2 + 1 + 4 = 7$$

↓

Implying:

$$\left\{ \begin{array}{l} x_1 = S \\ x_2 = T \end{array} \right.$$

Q: Which cut is better?

A: Cut 2 is cheaper than cut 1!

Q: How to find the minimum cut for all possible cuts?

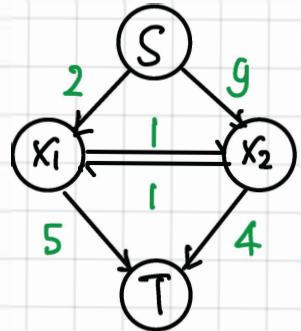
[5] Graphs: Augmenting Path

#4

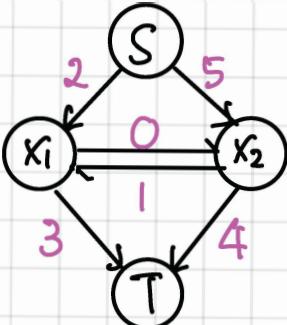
(the Ford - Fulkerson method)

Basic components

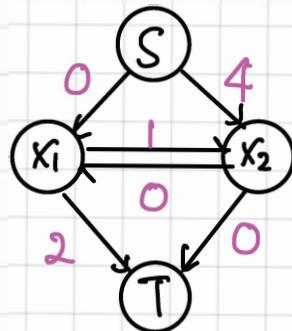
(1) Capacity graph



(2) Flow graph

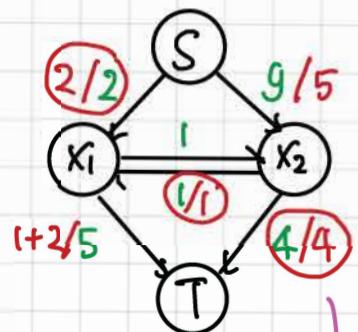


(3) Residual graph

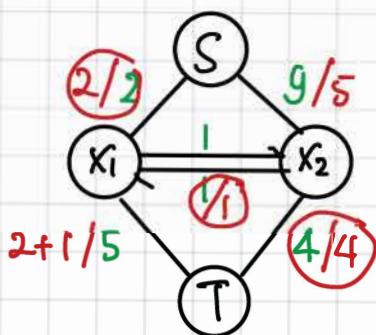


Calculation:

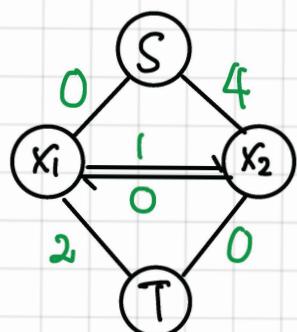
Start from 2 :



Alternative 1 (Start from 9):



Residual:

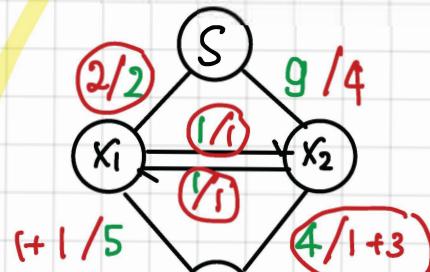


Must be the same!

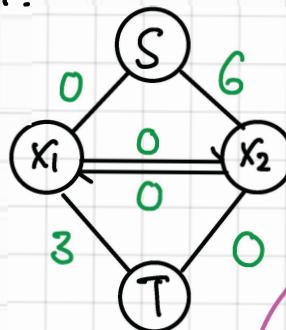
$$\text{Cut cost} = 2 + 1 + 4 = 7 \\ (\text{Min-cut})$$

$$\text{Max-flow} = 2 + 5 = 3 + 4 = 7$$

Alternative 2
(from 2 then 1):



Residual:

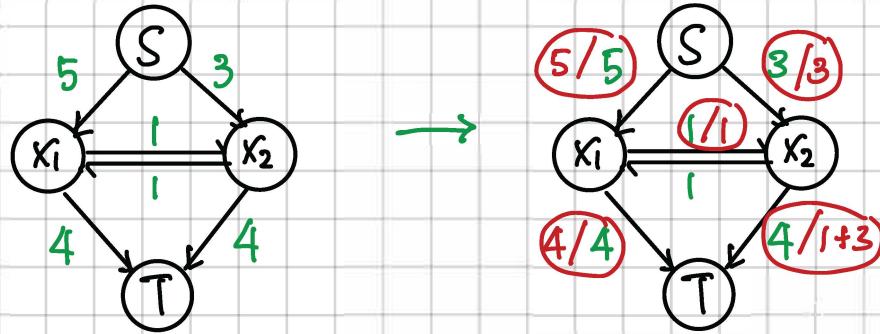


$$\text{Cut cost} = 2 + 1 + 1 + 4 = 8 \\ (\text{Min-cut})$$

$$\text{Max-flow} = 2 + 4 = 1 + 1 + 1 + 3 = 6$$

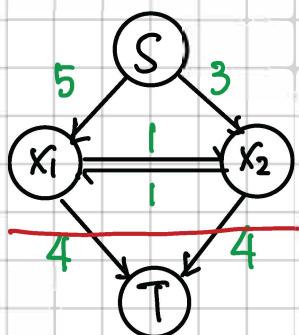
Other examples:

1.



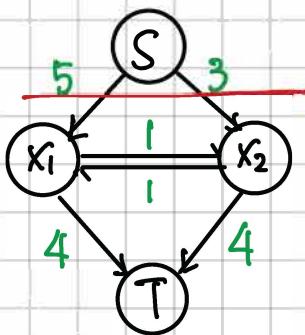
Min cut = 8
Maxflow = 8

For this case, there are a few different ways of cutting:



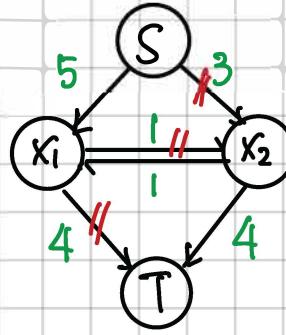
cut cost = 8

$x_1 = x_2 = T$



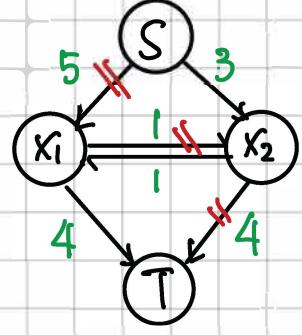
cut cost = 8

$x_1 = x_2 = S$



cut cost = 8

$x_1 = T; x_2 = S$



cut cost = 10

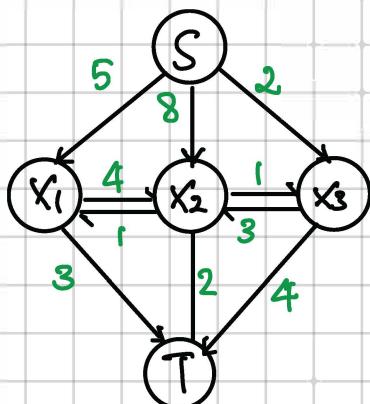
$x_1 = S, x_2 = T$

(not preferred because the cost is high)

all these labelings are correct.

This implies there are ambiguities in the labeling.

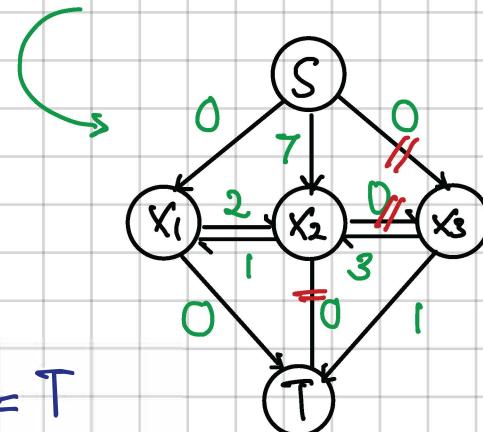
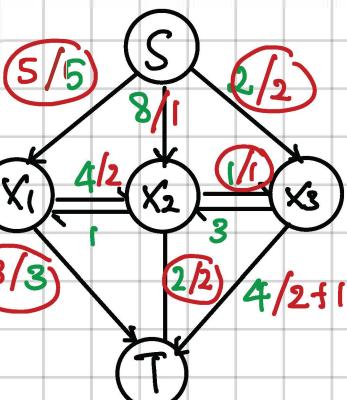
2.



$x_2 = T, x_3 = S \rightarrow \text{total cost} = 2+1+2 = 5$

$x_1 = \begin{cases} S & \rightarrow \text{additional cost} = 5 \\ T & \rightarrow \text{additional cost} = 3 \end{cases}$

We know the Max-flow = 8, thus $x_1 = T$



[6] Belief Propagation (BP)

Aside from graphcuts, BP is another popular method for graph optimization. It's more general than graphcuts (can accept many labels) but much slower.

Formulation:

→ to prevent a very small / large value due to the product of m_{ji}

Scalar
Data term

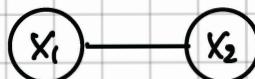
\downarrow \downarrow

$$\text{Belief: } b_i(x_i) = k \phi_i(x_i) \prod_{j \in N_i} m_{ji}(x_i)$$

$$\text{Message passing: } m_{ij}(x_j) = \sum_{x_i} \phi_i(x_i) \psi_{ij}(x_i, x_j) \prod_{k \in N_i \setminus j} m_{ki}(x_i)$$

↑
Prior term
↓
Neighbors of i except j.

Example:



$$; x_i = \{F, B\}$$

$$b(x_i=F) = k f_d(x_i=F, d) m_{2i}(x_i=F)$$

$$m_{2i}(x_i=F) = \sum_{x_2 \in \{F, B\}} f_d(x_2, d_2) f_p(x_2, x_i=F)$$

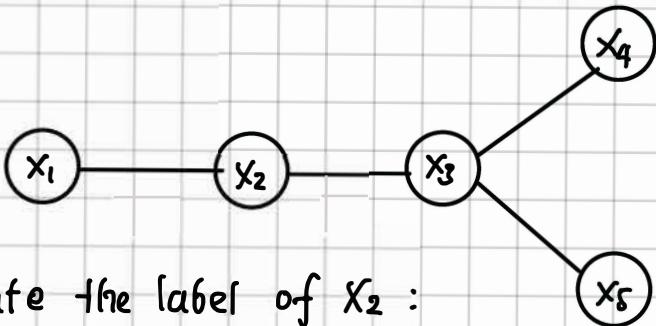
$$= f_d(x_2=F, d_2) f_p(x_2=F, x_i=F) + f_d(x_2=B, d_2) f_p(x_2=B, x_i=F)$$

Aside from $b(x_i=F)$, we also compute $b(x_i=B)$, and compare $b(x_i=F)$ and $b(x_i=B)$.

If $b(x_i=F) > b(x_i=B)$, then we label $x_i=B$

For minimization

Another example :



Suppose we want to calculate the label of x_2 :

$$b_2(x_2) = k \phi_2(x_2) m_{12}(x_2) m_{32}(x_2)$$

$$m_{12}(x_2) = \sum_{x_1} \phi_1(x_1) \psi_{12}(x_1, x_2)$$

$$m_{32}(x_2) = \sum_{x_3} \phi_3(x_3) \psi_{32}(x_3, x_2) m_{43}(x_3) m_{53}(x_3)$$

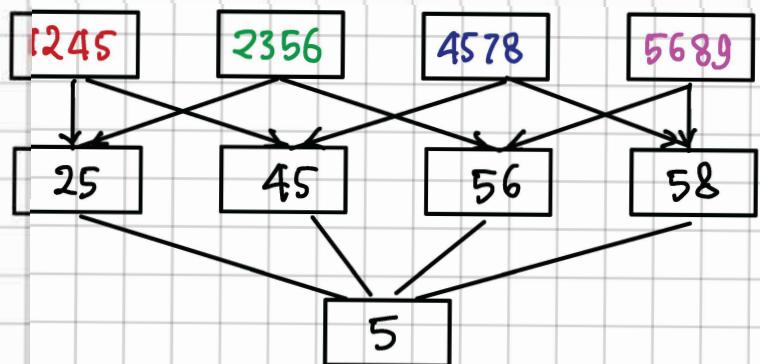
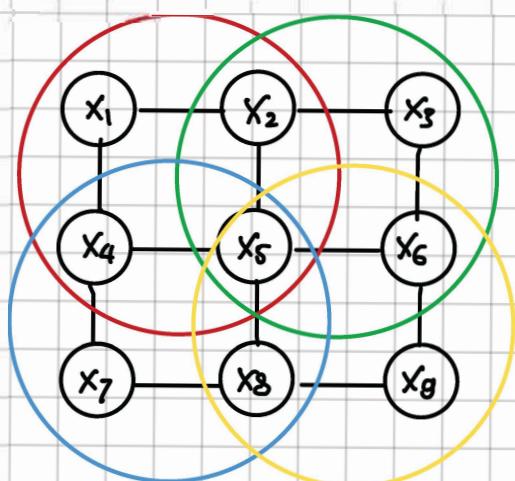
$$m_{43}(x_3) = \sum_{x_4} \phi_4(x_4) \psi_{43}(x_4, x_3)$$

$$m_{53}(x_3) = \sum_{x_5} \phi_5(x_5) \psi_{53}(x_5, x_3)$$

Note:

» The standard BP works only for a non-loopy graph.

For loopy graphs, one of the solutions is the Generalized BP.



$$b_5(x_5) = k \phi_5(x_5) m_{2 \rightarrow 5} m_{4 \rightarrow 5} m_{6 \rightarrow 5} m_{8 \rightarrow 5}$$

$$m_{4 \rightarrow 5}(x_5) = k \sum_{x_4} \phi_4(x_4) \psi_{45}(x_4, x_5) m_{12 \rightarrow 45} m_{78 \rightarrow 45}$$